

Assignment 5: Build and Evaluate Classification Models

Student: Abigail Albury-Bloom Focus: Multi-class obesity risk classification

1. Data loading and basic structure

```
In [1]: import pandas as pd, numpy as np
from pathlib import Path

# Expect train.csv and test.csv in the same folder as this notebook
train = pd.read_csv(Path('train.csv'))
test = pd.read_csv(Path('test.csv'))
target = 'NObeyesdad'
X = train.drop(columns=[target, 'id'])
y = train[target]

print('Train shape:', train.shape)
print('Test shape:', test.shape)
print('Target classes:', y.unique())
```

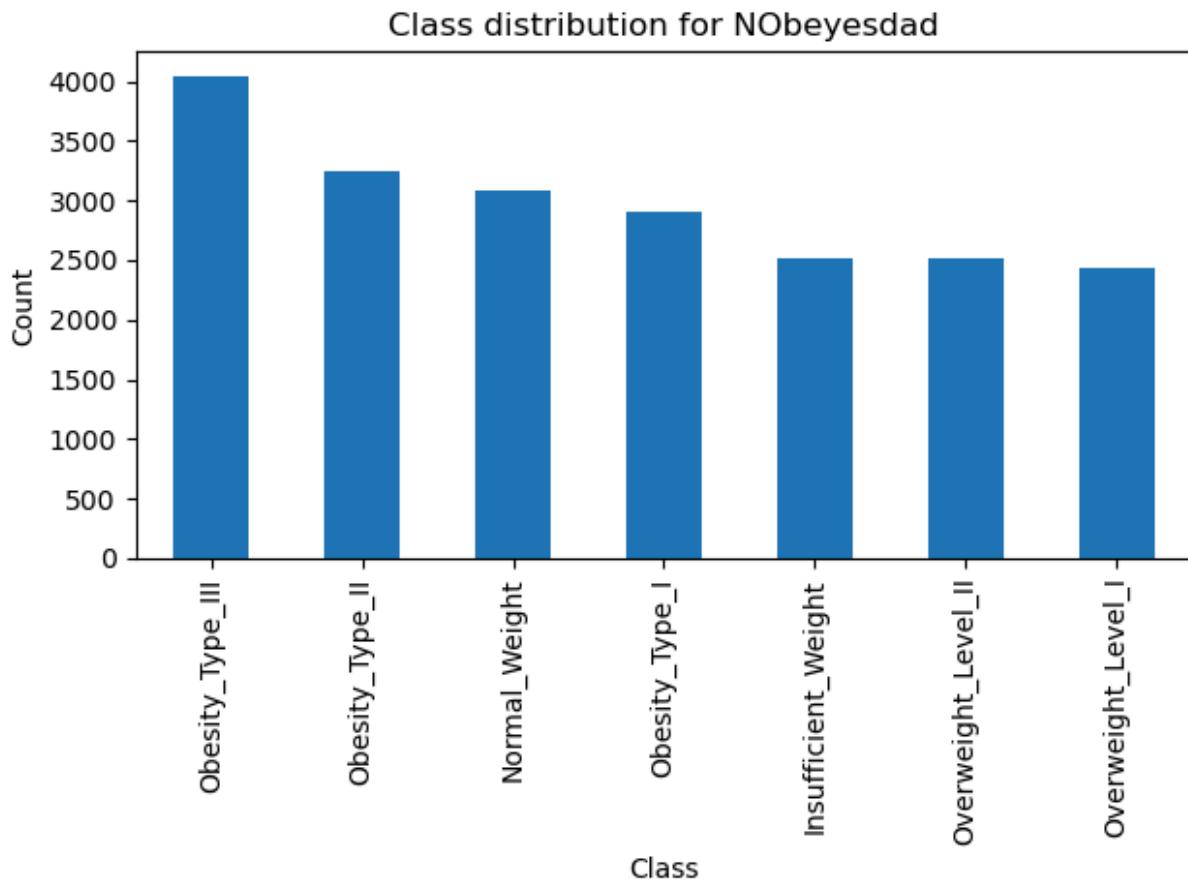
```
Train shape: (20758, 18)
Test shape: (13840, 17)
Target classes: ['Overweight_Level_II' 'Normal_Weight' 'Insufficient_Weight'
 'Obesity_Type_III' 'Obesity_Type_II' 'Overweight_Level_I'
 'Obesity_Type_I']
```

2. Class distribution and feature summary

```
In [2]: import matplotlib.pyplot as plt

plt.figure()
y.value_counts().plot(kind='bar')
plt.title('Class distribution for NObeyesdad')
plt.xlabel('Class')
plt.ylabel('Count')
plt.tight_layout()
plt.show()

print('Numeric feature summary:')
display(X.select_dtypes(include=[np.number]).describe().T)
```



Numeric feature summary:

	count	mean	std	min	25%	50%	75%	
Age	20758.0	23.841804	5.688072	14.00	20.000000	22.815416	26.000000	61.0
Height	20758.0	1.700245	0.087312	1.45	1.631856	1.700000	1.762887	1.9
Weight	20758.0	87.887768	26.379443	39.00	66.000000	84.064875	111.600553	165.0
FCVC	20758.0	2.445908	0.533218	1.00	2.000000	2.393837	3.000000	3.0
NCP	20758.0	2.761332	0.705375	1.00	3.000000	3.000000	3.000000	4.0
CH2O	20758.0	2.029418	0.608467	1.00	1.792022	2.000000	2.549617	3.0
FAF	20758.0	0.981747	0.838302	0.00	0.008013	1.000000	1.587406	3.0
TUE	20758.0	0.616756	0.602113	0.00	0.000000	0.573887	1.000000	2.0

3. Feature types and preprocessing pipeline

```
In [3]: from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler

numeric_cols = X.select_dtypes(include=[np.number]).columns.tolist()
categorical_cols = X.select_dtypes(exclude=[np.number]).columns.tolist()
print('Numeric columns:', numeric_cols)
print('Categorical columns:', categorical_cols)
```

```
preprocess = ColumnTransformer([
    ('num', StandardScaler(), numeric_cols),
    ('cat', OneHotEncoder(handle_unknown='ignore', sparse_output=False), cat_cols)
])
```

Numeric columns: ['Age', 'Height', 'Weight', 'FCVC', 'NCP', 'CH20', 'FAF', 'TUE']

Categorical columns: ['Gender', 'family_history_with_overweight', 'FAVC', 'CAEC', 'SMOKE', 'SCC', 'CALC', 'MTRANS']

4. Cross-validated performance for four classifiers

```
In [4]: from sklearn.pipeline import Pipeline
from sklearn.model_selection import StratifiedKFold, cross_val_score
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

base_models = {
    'Logistic_L2': LogisticRegression(penalty='l2', max_iter=1000),
    'LDA': LinearDiscriminantAnalysis(),
    'GaussianNB': GaussianNB(),
    'SVM_RBF': SVC(kernel='rbf', C=1, gamma='scale', decision_function_shape='ovr')
}

cv_results = []
for name, clf in base_models.items():
    pipe = Pipeline([('pre', preprocess), ('clf', clf)])
    scores = cross_val_score(pipe, X, y, cv=cv, scoring='accuracy')
    cv_results.append({'Model': name, 'Mean_Accuracy': scores.mean(), 'Std_Acc': scores.std()})
    print(f'{name}: {scores.mean():.4f} mean ± {scores.std():.4f} sd")
```

Logistic_L2: 0.8620 mean ± 0.0046 sd

LDA: 0.8207 mean ± 0.0032 sd

GaussianNB: 0.5874 mean ± 0.0053 sd

SVM_RBF: 0.8788 mean ± 0.0058 sd

5. Cross-validated accuracy comparison table

```
In [5]: cv_df = pd.DataFrame(cv_results).sort_values(by='Mean_Accuracy', ascending=False)
display(cv_df)
```

	Model	Mean_Accuracy	Std_Accuracy
0	SVM_RBF	0.878841	0.005776
1	Logistic_L2	0.862029	0.004608
2	LDA	0.820744	0.003238
3	GaussianNB	0.587388	0.005335

6. Holdout validation and confusion matrices

```
In [6]: from sklearn.model_selection import train_test_split

X_train, X_val, y_train, y_val = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)

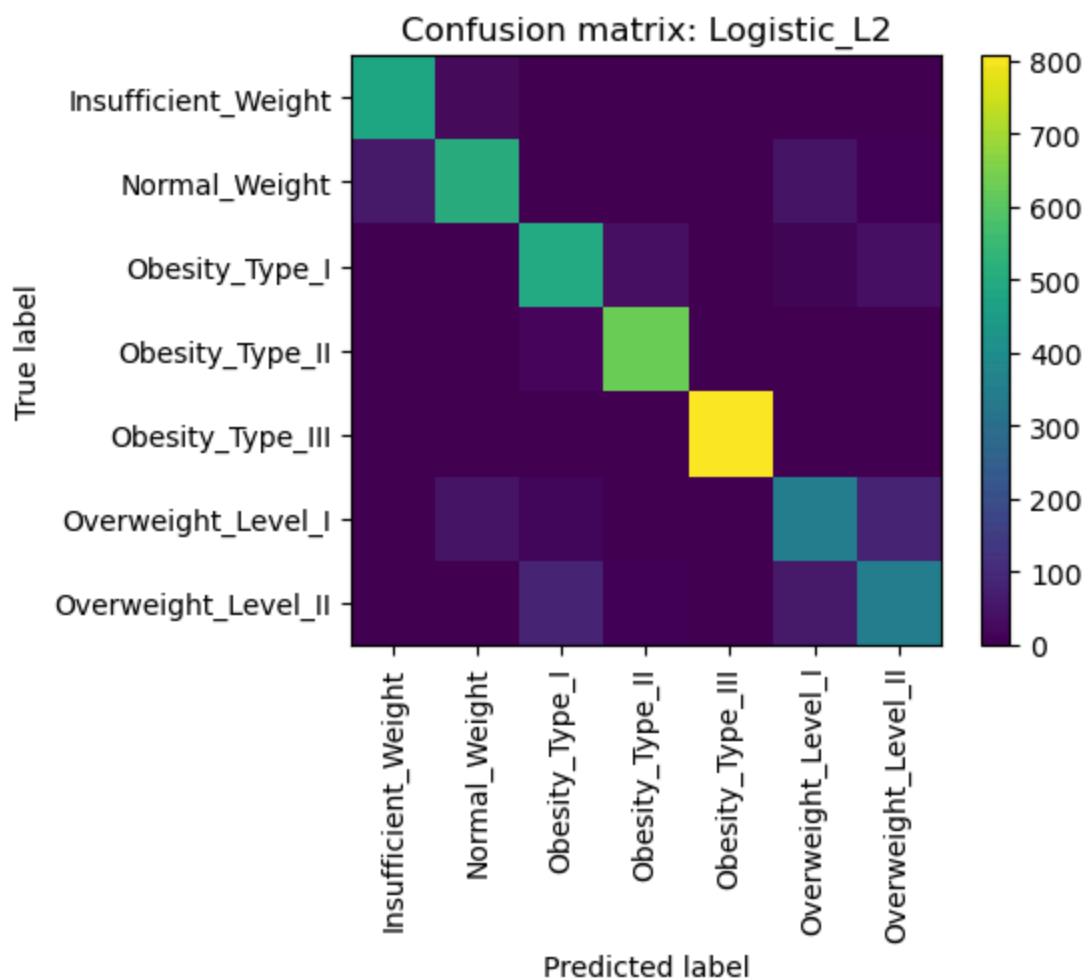
validation_results = []
label_order = sorted(y.unique())

for name, clf in base_models.items():
    pipe = Pipeline([('pre', preprocess), ('clf', clf)])
    pipe.fit(X_train, y_train)
    preds = pipe.predict(X_val)
    acc = accuracy_score(y_val, preds)
    validation_results.append({'Model': name, 'Holdout_Accuracy': acc})
    print(f"\n{name} holdout accuracy: {acc:.4f}")
    cm = confusion_matrix(y_val, preds, labels=label_order)
    cm_df = pd.DataFrame(cm, index=label_order, columns=label_order)
    display(cm_df)

# Confusion matrix heatmap for visuals
plt.figure(figsize=(6, 5))
plt.imshow(cm, interpolation='nearest')
plt.title(f'Confusion matrix: {name}')
plt.colorbar()
tick_marks = range(len(label_order))
plt.xticks(tick_marks, label_order, rotation=90)
plt.yticks(tick_marks, label_order)
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.tight_layout()
plt.show()
```

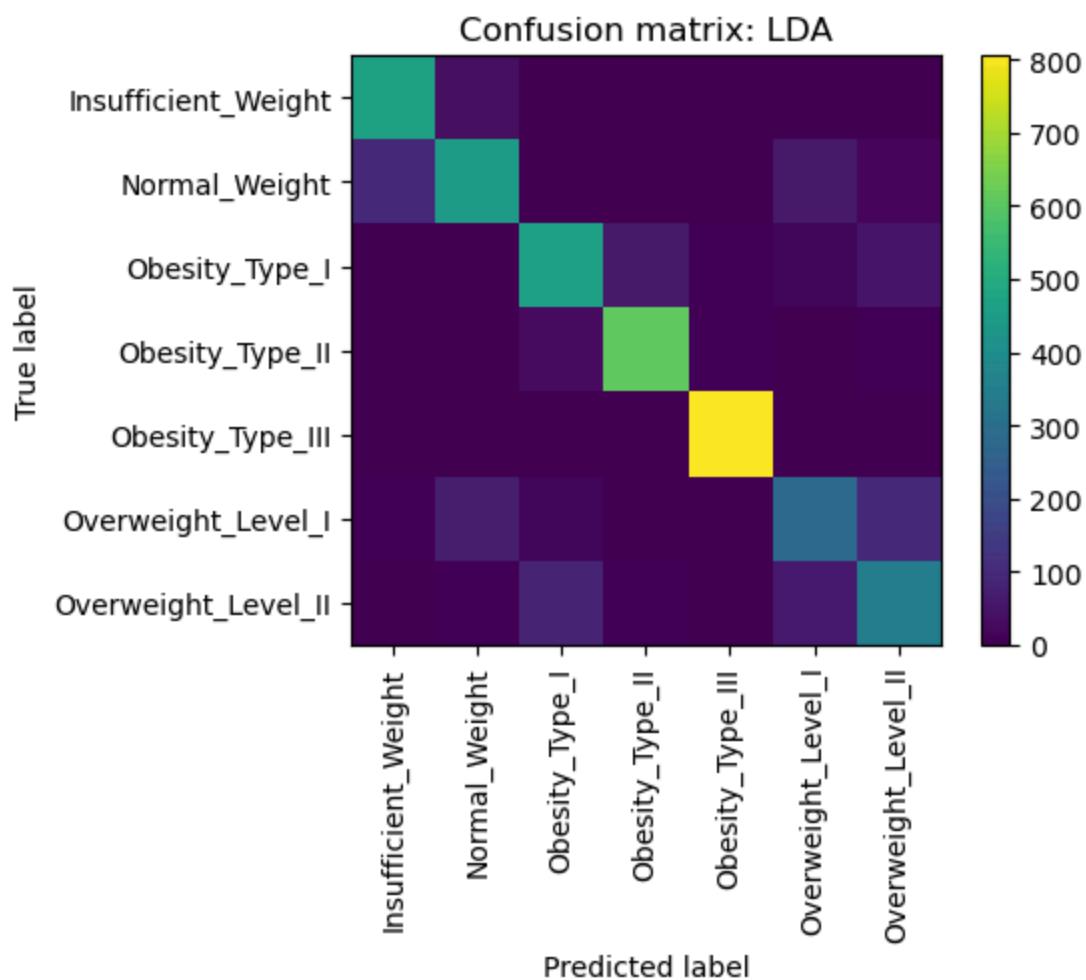
Logistic_L2 holdout accuracy: 0.8683

	Insufficient_Weight	Normal_Weight	Obesity_Type_I	Obesity_Type_
Insufficient_Weight	479	25	0	
Normal_Weight	58	505	2	
Obesity_Type_I	1	0	492	3
Obesity_Type_II	0	0	22	62
Obesity_Type_III	0	0	0	
Overweight_Level_I	1	46	15	
Overweight_Level_II	0	3	79	



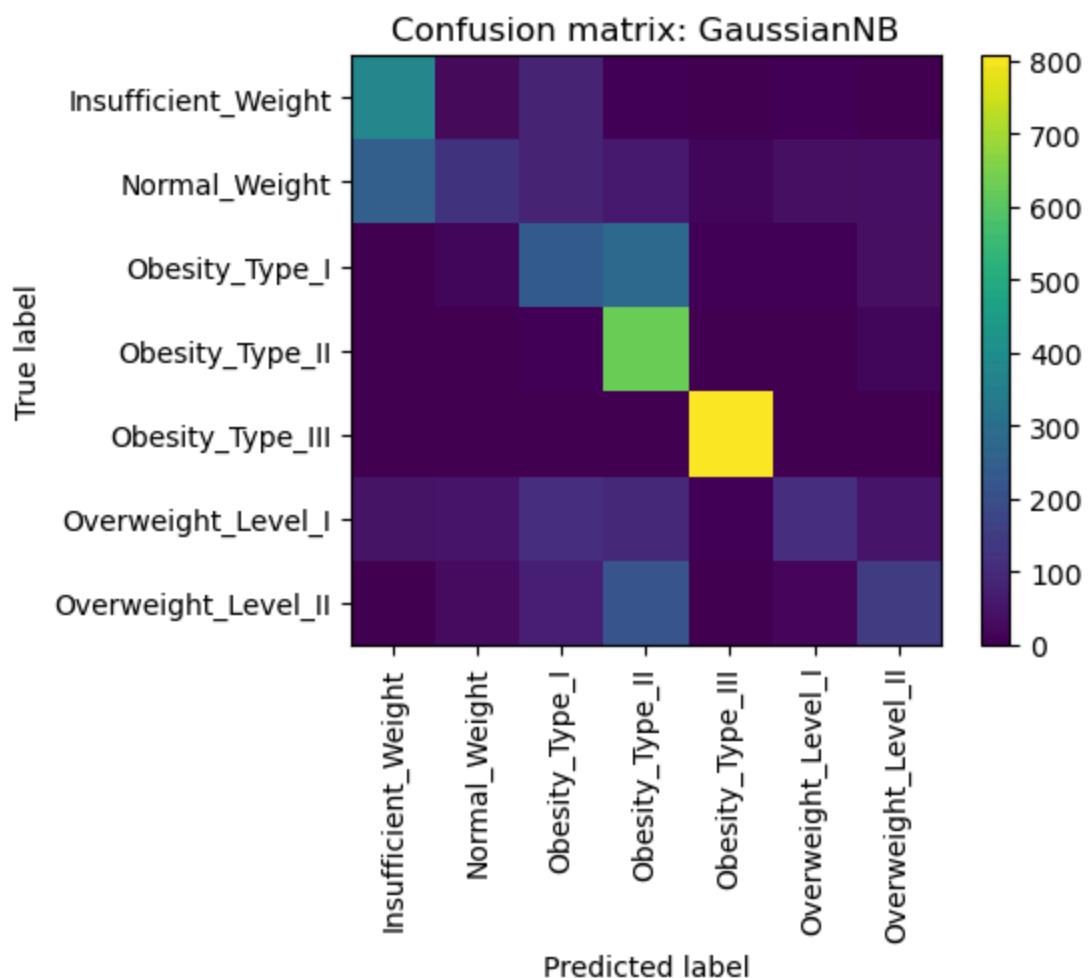
LDA holdout accuracy: 0.8230

	Insufficient_Weight	Normal_Weight	Obesity_Type_I	Obesity_Type_II
Insufficient_Weight	470	33	0	5
Normal_Weight	101	442	1	6
Obesity_Type_I	1	1	456	5
Obesity_Type_II	0	0	29	6
Obesity_Type_III	0	0	1	1
Overweight_Level_I	6	74	16	5
Overweight_Level_II	0	9	82	6



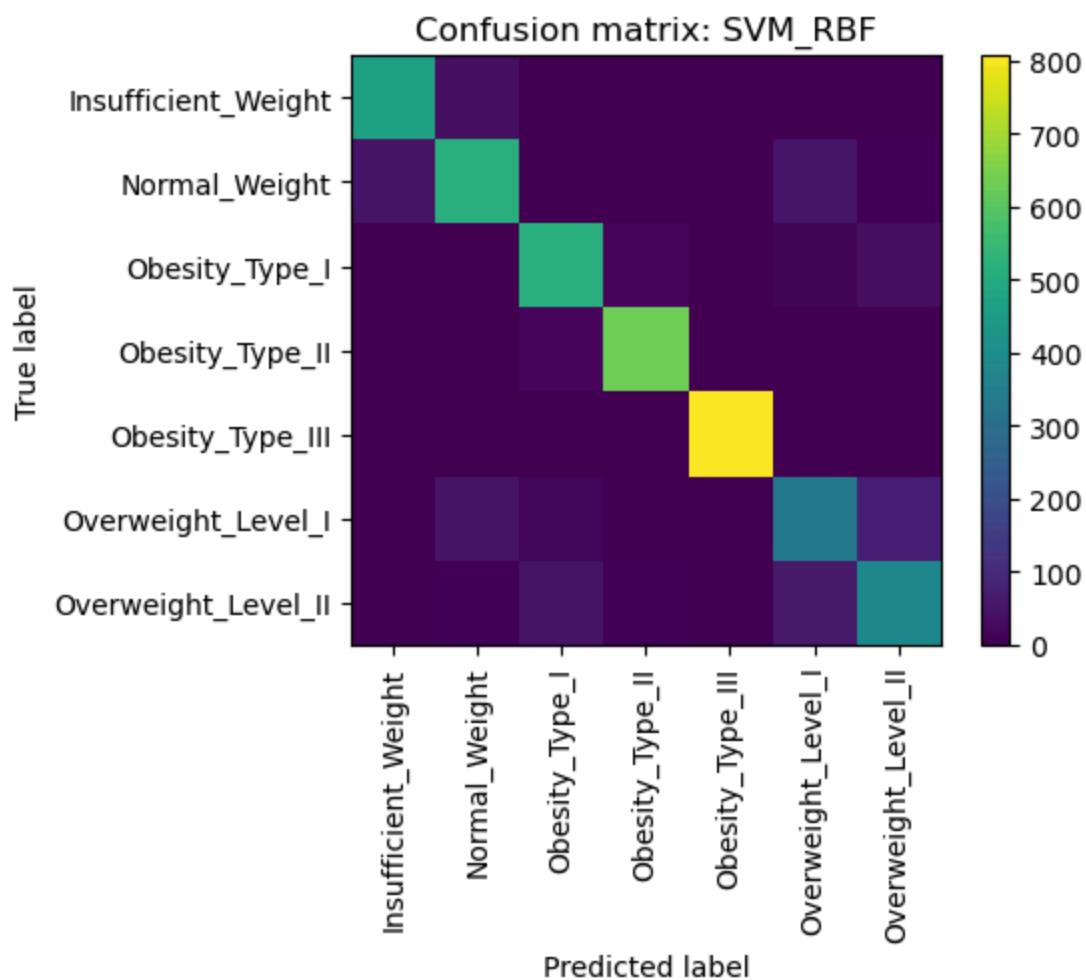
GaussianNB holdout accuracy: 0.5860

	Insufficient_Weight	Normal_Weight	Obesity_Type_I	Obesity_Type_II
Insufficient_Weight	377	25	89	6
Normal_Weight	257	123	82	6
Obesity_Type_I	1	17	239	27
Obesity_Type_II	0	2	5	62
Obesity_Type_III	0	1	0	0
Overweight_Level_I	42	53	113	10
Overweight_Level_II	2	29	77	22



SVM_RBF holdout accuracy: 0.8810

	Insufficient_Weight	Normal_Weight	Obesity_Type_I	Obesity_Type_II
Insufficient_Weight	472	32	0	0
Normal_Weight	47	510	1	0
Obesity_Type_I	1	0	513	0
Obesity_Type_II	0	0	19	63
Obesity_Type_III	0	0	2	0
Overweight_Level_I	2	47	18	0
Overweight_Level_II	0	6	47	0



7. Holdout accuracy comparison table

```
In [7]: holdout_df = pd.DataFrame(validation_results).sort_values(by='Holdout_Accuracy')
display(holdout_df)
```

	Model	Holdout_Accuracy
0	SVM_RBF	0.881021
1	Logistic_L2	0.868256
2	LDA	0.822977
3	GaussianNB	0.585983

8. Train final models on full data and generate prediction files

```
In [8]: final_predictions = {}
X_test_features = test.drop(columns=['id'])

for name, clf in base_models.items():
    pipe = Pipeline([('pre', preprocess), ('clf', clf)])
    pipe.fit(X, y)
    preds = pipe.predict(X_test_features)
    final_predictions[name] = preds
    out_df = pd.DataFrame({'id': test['id'], target: preds})
    out_name = f'submission_{name}.csv'
    out_df.to_csv(out_name, index=False)
    print(f'Saved {out_name}')


Saved submission_Logistic_L2.csv
Saved submission_LDA.csv
Saved submission_GaussianNB.csv
Saved submission_SVM_RBF.csv
```

1. Data loading and initial structure

```
In [9]: import pandas as pd
import numpy as np
from pathlib import Path

# Expect train.csv, test.csv, and sample_submission.csv to be in the same folder
base_path = Path('.')
train_path = base_path / 'train.csv'
test_path = base_path / 'test.csv'
sample_path = base_path / 'sample_submission.csv'

train = pd.read_csv(train_path)
test = pd.read_csv(test_path)
sample = pd.read_csv(sample_path)

target = 'NObeyesdad'
X = train.drop(columns=[target, 'id'])
y = train[target]

X.head()
```

Out [9]:

	Gender	Age	Height	Weight	family_history_with_overweight	FAVC	
0	Male	24.443011	1.699998	81.669950		yes	yes 2.
1	Female	18.000000	1.560000	57.000000		yes	yes 2.
2	Female	18.000000	1.711460	50.165754		yes	yes 1.
3	Female	20.952737	1.710730	131.274851		yes	yes 3.
4	Male	31.641081	1.914186	93.798055		yes	yes 2.

2. Feature types and preprocessing pipeline

In [10]:

```
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.pipeline import Pipeline

numeric_cols = X.select_dtypes(include=[np.number]).columns.tolist()
categorical_cols = X.select_dtypes(exclude=[np.number]).columns.tolist()
print('Numeric columns:', numeric_cols)
print('Categorical columns:', categorical_cols)

preprocess = ColumnTransformer([
    ('num', StandardScaler(), numeric_cols),
    ('cat', OneHotEncoder(handle_unknown='ignore', sparse_output=False), cat)])

```

Numeric columns: ['Age', 'Height', 'Weight', 'FCVC', 'NCP', 'CH20', 'FAF', 'TUE']
Categorical columns: ['Gender', 'family_history_with_overweight', 'FAVC', 'C AEC', 'SMOKE', 'SCC', 'CALC', 'MTRANS']

3. Class distribution and numeric correlations

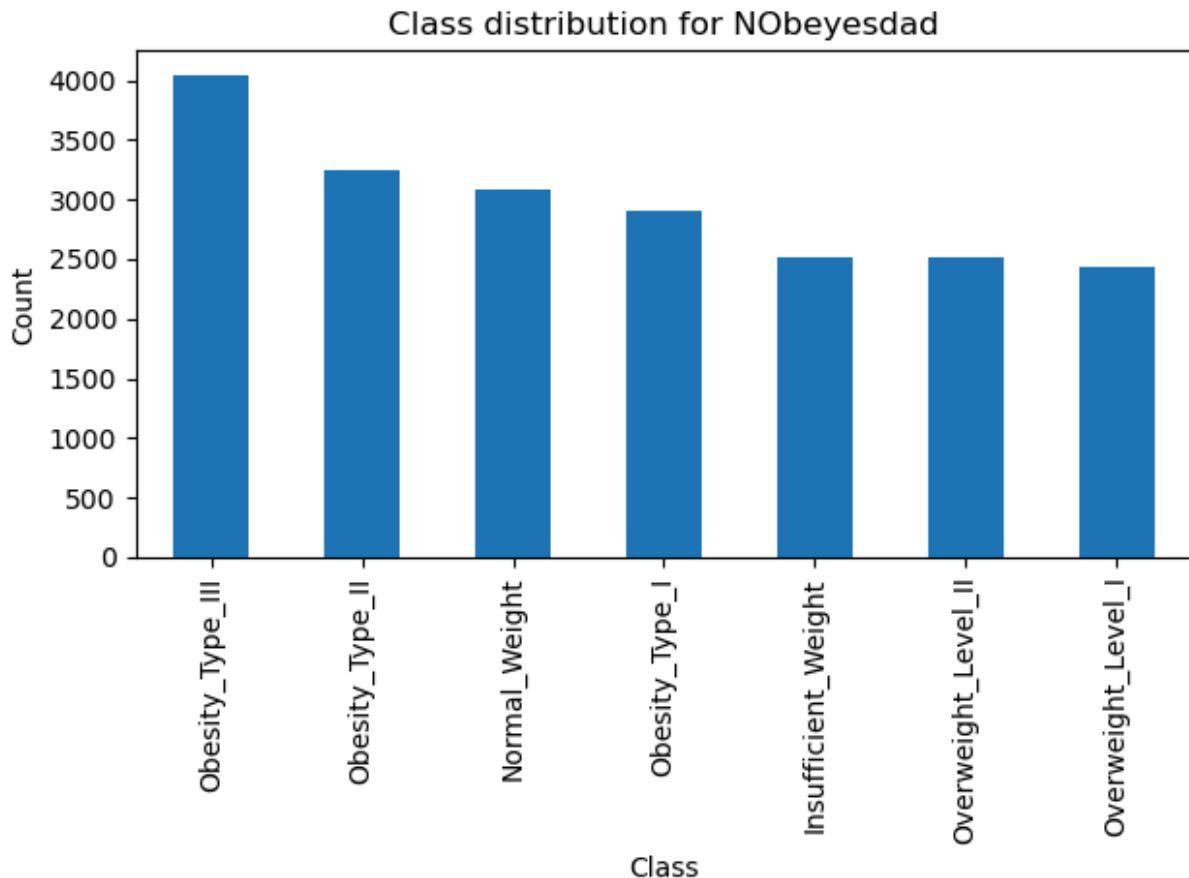
In [11]:

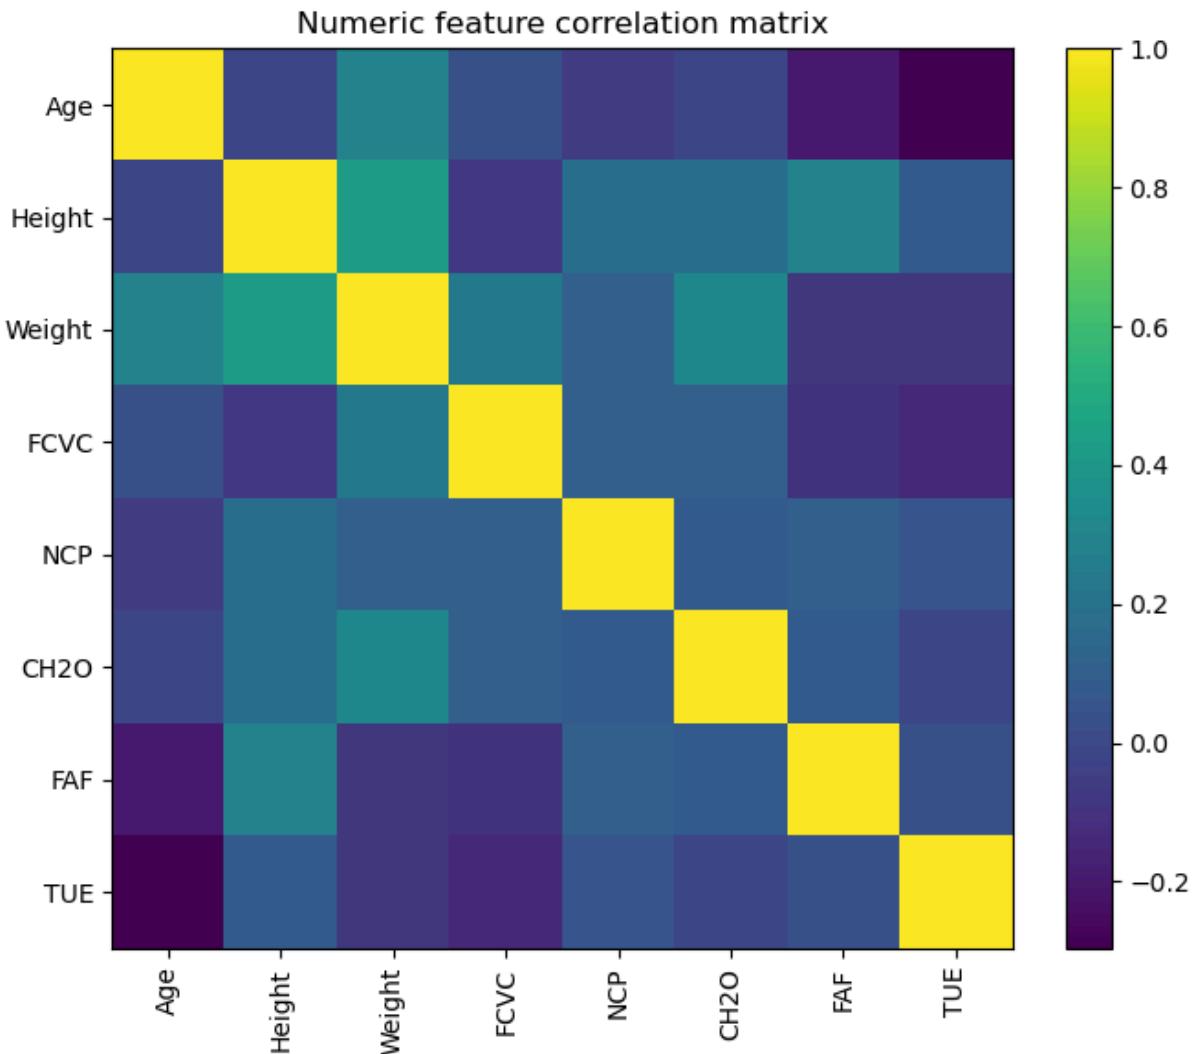
```
import matplotlib.pyplot as plt

# Class distribution
plt.figure()
train[target].value_counts().plot(kind='bar')
plt.title('Class distribution for NOobeyesdad')
plt.xlabel('Class')
plt.ylabel('Count')
plt.tight_layout()
plt.show()

# Correlation heatmap for numeric features
corr = train[numeric_cols].corr()
plt.figure(figsize=(8, 6))
im = plt.imshow(corr, interpolation='nearest')
plt.title('Numeric feature correlation matrix')
plt.colorbar(im, fraction=0.046, pad=0.04)
plt.xticks(range(len(numeric_cols)), numeric_cols, rotation=90)
plt.yticks(range(len(numeric_cols)), numeric_cols)
```

```
plt.tight_layout()  
plt.show()
```





4. Cross-validated model performance (four classifiers)

```
In [12]: from sklearn.model_selection import StratifiedKFold, cross_val_score
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

def train_and_eval(name, model):
    pipe = Pipeline([('pre', preprocess), ('clf', model)])
    scores = cross_val_score(pipe, X, y, cv=cv, scoring='accuracy')
    print(f"{name}: {scores.mean():.4f} mean ± {scores.std():.4f} sd")
    pipe.fit(X, y)
    return pipe

# 1. Multinomial logistic regression with L2 regularization
logit_model = train_and_eval('Multinomial logistic regression (L2)', Logisti
```

```

# 2. Linear Discriminant Analysis
lda_model = train_and_eval('Linear Discriminant Analysis', LinearDiscriminatoryAnalysis)

# 3. Gaussian Naive Bayes
nb_model = train_and_eval('Gaussian Naive Bayes', GaussianNB())

# 4. Support Vector Machine with RBF kernel
svm_model = train_and_eval('RBF-kernel Support Vector Machine', SVC(kernel='rbf'))

```

Multinomial logistic regression (L2): 0.8620 mean ± 0.0046 sd
 Linear Discriminant Analysis: 0.8207 mean ± 0.0032 sd
 Gaussian Naive Bayes: 0.5874 mean ± 0.0053 sd
 RBF-kernel Support Vector Machine: 0.8788 mean ± 0.0058 sd

5. Validation split and confusion matrices

```
In [13]: from sklearn.model_selection import train_test_split

X_train, X_val, y_train, y_val = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)

models_for_validation = [
    ('Multinomial logistic regression (L2)', LogisticRegression(penalty='l2')),
    ('Linear Discriminant Analysis', LinearDiscriminantAnalysis()),
    ('Gaussian Naive Bayes', GaussianNB()),
    ('RBF-kernel Support Vector Machine', SVC(kernel='rbf', C=1, gamma='scale'))
]

for name, clf in models_for_validation:
    pipe = Pipeline([('pre', preprocess), ('clf', clf)])
    pipe.fit(X_train, y_train)
    preds = pipe.predict(X_val)
    acc = accuracy_score(y_val, preds)
    print(f"\n{name} validation accuracy: {acc:.4f}")
    cm = confusion_matrix(y_val, preds, labels=sorted(y.unique()))
    cm_df = pd.DataFrame(cm, index=sorted(y.unique()), columns=sorted(y.unique()))
    print(cm_df)
```

Multinomial logistic regression (L2) validation accuracy: 0.8683

	Insufficient_Weight	Normal_Weight	Obesity_Type_I	\
Insufficient_Weight	479	25	0	
Normal_Weight	58	505	2	
Obesity_Type_I	1	0	492	
Obesity_Type_II	0	0	22	
Obesity_Type_III	0	0	0	
Overweight_Level_I	1	46	15	
Overweight_Level_II	0	3	79	
	Obesity_Type_II	Obesity_Type_III	Overweight_Level_I	
\				
Insufficient_Weight	0	0	1	
Normal_Weight	0	0	43	
Obesity_Type_I	37	3	11	
Obesity_Type_II	625	0	0	
Obesity_Type_III	1	807	1	
Overweight_Level_I	0	0	342	
Overweight_Level_II	8	0	59	
	Overweight_Level_II			
Insufficient_Weight	0			
Normal_Weight	9			
Obesity_Type_I	38			
Obesity_Type_II	3			
Obesity_Type_III	0			
Overweight_Level_I	81			
Overweight_Level_II	355			

Linear Discriminant Analysis validation accuracy: 0.8230

	Insufficient_Weight	Normal_Weight	Obesity_Type_I	\
Insufficient_Weight	470	33	0	
Normal_Weight	101	442	1	
Obesity_Type_I	1	1	456	
Obesity_Type_II	0	0	29	
Obesity_Type_III	0	0	1	
Overweight_Level_I	6	74	16	
Overweight_Level_II	0	9	82	
	Obesity_Type_II	Obesity_Type_III	Overweight_Level_I	
\				
Insufficient_Weight	0	0	1	
Normal_Weight	0	0	54	
Obesity_Type_I	54	7	13	
Obesity_Type_II	612	5	0	
Obesity_Type_III	2	805	1	
Overweight_Level_I	0	0	289	
Overweight_Level_II	4	0	66	
	Overweight_Level_II			
Insufficient_Weight	1			
Normal_Weight	19			
Obesity_Type_I	50			
Obesity_Type_II	4			
Obesity_Type_III	0			
Overweight_Level_I	100			

Overweight_Level_II	343			
Gaussian Naive Bayes validation accuracy: 0.5860				
	Insufficient_Weight	Normal_Weight	Obesity_Type_I	\
Insufficient_Weight	377	25	89	
Normal_Weight	257	123	82	
Obesity_Type_I	1	17	239	
Obesity_Type_II	0	2	5	
Obesity_Type_III	0	1	0	
Overweight_Level_I	42	53	113	
Overweight_Level_II	2	29	77	
	Obesity_Type_II	Obesity_Type_III	Overweight_Level_I	
\				
Insufficient_Weight	4	2	5	
Normal_Weight	61	17	37	
Obesity_Type_I	275	5	9	
Obesity_Type_II	624	0	3	
Obesity_Type_III	1	806	1	
Overweight_Level_I	103	9	114	
Overweight_Level_II	223	1	22	
	Overweight_Level_II			
Insufficient_Weight	3			
Normal_Weight	40			
Obesity_Type_I	36			
Obesity_Type_II	16			
Obesity_Type_III	0			
Overweight_Level_I	51			
Overweight_Level_II	150			
RBF-kernel Support Vector Machine validation accuracy: 0.8810				
	Insufficient_Weight	Normal_Weight	Obesity_Type_I	\
Insufficient_Weight	472	32	0	
Normal_Weight	47	510	1	
Obesity_Type_I	1	0	513	
Obesity_Type_II	0	0	19	
Obesity_Type_III	0	0	2	
Overweight_Level_I	2	47	18	
Overweight_Level_II	0	6	47	
	Obesity_Type_II	Obesity_Type_III	Overweight_Level_I	
\				
Insufficient_Weight	0	0	1	
Normal_Weight	0	0	51	
Obesity_Type_I	21	3	12	
Obesity_Type_II	630	0	0	
Obesity_Type_III	1	806	0	
Overweight_Level_I	1	0	340	
Overweight_Level_II	6	0	58	
	Overweight_Level_II			
Insufficient_Weight	0			
Normal_Weight	8			
Obesity_Type_I	32			
Obesity_Type_II	1			

Obesity_Type_III	0
Overweight_Level_I	77
Overweight_Level_II	387

6. Generate prediction files for all four models

```
In [14]: # Fit each model on the full training data and generate prediction files
logit_pipe = Pipeline([('pre', preprocess), ('clf', LogisticRegression(penalty='l2'))])
lda_pipe = Pipeline([('pre', preprocess), ('clf', LinearDiscriminantAnalysis())])
nb_pipe = Pipeline([('pre', preprocess), ('clf', GaussianNB())])
svm_pipe = Pipeline([('pre', preprocess), ('clf', SVC(kernel='rbf', C=1, gamma=0.001))])

logit_pipe.fit(X, y)
lda_pipe.fit(X, y)
nb_pipe.fit(X, y)
svm_pipe.fit(X, y)

logit_preds = logit_pipe.predict(test.drop(columns=['id']))
lda_preds = lda_pipe.predict(test.drop(columns=['id']))
nb_preds = nb_pipe.predict(test.drop(columns=['id']))
svm_preds = svm_pipe.predict(test.drop(columns=['id']))

pd.DataFrame({'id': test['id'], target: logit_preds}).to_csv('submission_logistic_l2.csv')
pd.DataFrame({'id': test['id'], target: lda_preds}).to_csv('submission_lda.csv')
pd.DataFrame({'id': test['id'], target: nb_preds}).to_csv('submission_gaussian_nb.csv')
pd.DataFrame({'id': test['id'], target: svm_preds}).to_csv('submission_svm_rbf.csv')

print('Prediction files generated:')
print(' - submission_logistic_l2.csv')
print(' - submission_lda.csv')
print(' - submission_gaussian_nb.csv')
print(' - submission_svm_rbf.csv')
```

Prediction files generated:
- submission_logistic_l2.csv
- submission_lda.csv
- submission_gaussian_nb.csv
- submission_svm_rbf.csv