

# Assignment 5: Build and Evaluate Classification Models

## 1. Data loading and basic structure

```
In [52]: import pandas as pd, numpy as np
from pathlib import Path

# Expect train.csv and test.csv in the same folder as this notebook
train = pd.read_csv(Path('train.csv'))
test = pd.read_csv(Path('test.csv'))
target = 'NObeyesad'
X = train.drop(columns=[target, 'id'])
y = train[target]

print('Train shape:', train.shape)
print('Test shape:', test.shape)
print('Target classes:', y.unique())
```

```
Train shape: (20758, 18)
Test shape: (13840, 17)
Target classes: ['Overweight_Level_II' 'Normal_Weight' 'Insufficient_Weight'
 'Obesity_Type_III' 'Obesity_Type_II' 'Overweight_Level_I'
 'Obesity_Type_I']
```

```
In [53]: ## Model Assumptions – Multicollinearity
from sklearn.preprocessing import StandardScaler
from statsmodels.stats.outliers_influence import variance_inflation_factor
import pandas as pd

# First, check what columns are actually available in your DataFrame
# Uncomment the next line to see available columns
# print(train.columns)

# Define numeric columns from your dataset
# Use actual column names from your DataFrame instead of placeholders
# For example, if your DataFrame has columns like 'age', 'income', 'years_ex
numeric_cols = train.select_dtypes(include=['float64', 'int64']).columns.tol
# Or manually specify columns that exist in your DataFrame:
# numeric_cols = ['age', 'income', 'years_experience']

# Scale numeric features for VIF calculation
scaler = StandardScaler()
X_num = scaler.fit_transform(train[numeric_cols])
X_vif = pd.DataFrame(X_num, columns=numeric_cols)

# Compute VIF
vif_data = pd.DataFrame({
    "Feature": X_vif.columns,
    "VIF": [variance_inflation_factor(X_vif.values, i) for i in range(X_vif.
})
```

```
vif_data = vif_data.sort_values(by="VIF", ascending=False)
vif_data
```

Out [53]:

	Feature	VIF
3	Weight	1.667178
2	Height	1.504342
1	Age	1.251766
7	FAF	1.202175
4	FCVC	1.155249
6	CH2O	1.141256
8	TUE	1.135968
5	NCP	1.066736
0	id	1.000723


## 2. Class distribution and feature summary

In [54]:

```
import matplotlib.pyplot as plt

plt.figure()
y.value_counts().plot(kind='bar')
plt.title('Class distribution for NObeyesdad')
plt.xlabel('Class')
plt.ylabel('Count')
plt.tight_layout()
plt.show()

print('Numeric feature summary:')
display(X.select_dtypes(include=[np.number]).describe().T)
```



Numeric feature summary:

	count	mean	std	min	25%	50%	75%	
<b>Age</b>	20758.0	23.841804	5.688072	14.00	20.000000	22.815416	26.000000	61.0
<b>Height</b>	20758.0	1.700245	0.087312	1.45	1.631856	1.700000	1.762887	1.9
<b>Weight</b>	20758.0	87.887768	26.379443	39.00	66.000000	84.064875	111.600553	165.0
<b>FCVC</b>	20758.0	2.445908	0.533218	1.00	2.000000	2.393837	3.000000	3.0
<b>NCP</b>	20758.0	2.761332	0.705375	1.00	3.000000	3.000000	3.000000	4.0
<b>CH2O</b>	20758.0	2.029418	0.608467	1.00	1.792022	2.000000	2.549617	3.0
<b>FAF</b>	20758.0	0.981747	0.838302	0.00	0.008013	1.000000	1.587406	3.0
<b>TUE</b>	20758.0	0.616756	0.602113	0.00	0.000000	0.573887	1.000000	2.0

### 3. Feature types and preprocessing pipeline

```
In [55]: from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler

numeric_cols = X.select_dtypes(include=[np.number]).columns.tolist()
categorical_cols = X.select_dtypes(exclude=[np.number]).columns.tolist()
print('Numeric columns:', numeric_cols)
print('Categorical columns:', categorical_cols)
```

```
preprocess = ColumnTransformer([
    ('num', StandardScaler(), numeric_cols),
    ('cat', OneHotEncoder(handle_unknown='ignore', sparse_output=False), cat_cols)
])
```

Numeric columns: ['Age', 'Height', 'Weight', 'FCVC', 'NCP', 'CH20', 'FAF', 'TUE']

Categorical columns: ['Gender', 'family\_history\_with\_overweight', 'FAVC', 'CAEC', 'SMOKE', 'SCC', 'CALC', 'MTRANS']

## 4. Cross-validated performance for four classifiers

```
In [56]: from sklearn.pipeline import Pipeline
from sklearn.model_selection import StratifiedKFold, cross_val_score
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

base_models = {
    'Logistic_L2': LogisticRegression(penalty='l2', max_iter=1000),
    'LDA': LinearDiscriminantAnalysis(),
    'GaussianNB': GaussianNB(),
    'SVM_RBF': SVC(kernel='rbf', C=1, gamma='scale', decision_function_shape='ovr')
}

cv_results = []
for name, clf in base_models.items():
    pipe = Pipeline([('pre', preprocess), ('clf', clf)])
    scores = cross_val_score(pipe, X, y, cv=cv, scoring='accuracy')
    cv_results.append({'Model': name, 'Mean_Accuracy': scores.mean(), 'Std_Acc': scores.std()})
    print(f'{name}: {scores.mean():.4f} mean ± {scores.std():.4f} sd")
```

Logistic\_L2: 0.8620 mean ± 0.0046 sd

LDA: 0.8207 mean ± 0.0032 sd

GaussianNB: 0.5874 mean ± 0.0053 sd

SVM\_RBF: 0.8788 mean ± 0.0058 sd

## 5. Cross-validated accuracy comparison table

```
In [57]: cv_df = pd.DataFrame(cv_results).sort_values(by='Mean_Accuracy', ascending=False)
display(cv_df)
```

	Model	Mean_Accuracy	Std_Accuracy
0	SVM_RBF	0.878841	0.005776
1	Logistic_L2	0.862029	0.004608
2	LDA	0.820744	0.003238
3	GaussianNB	0.587388	0.005335

## 6. Holdout validation and confusion matrices

```
In [58]: from sklearn.model_selection import train_test_split

X_train, X_val, y_train, y_val = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)


validation_results = []
label_order = sorted(y.unique())

for name, clf in base_models.items():
    pipe = Pipeline([('pre', preprocess), ('clf', clf)])
    pipe.fit(X_train, y_train)
    preds = pipe.predict(X_val)
    acc = accuracy_score(y_val, preds)
    validation_results.append({'Model': name, 'Holdout_Accuracy': acc})
    print(f"\n{name} holdout accuracy: {acc:.4f}")
    cm = confusion_matrix(y_val, preds, labels=label_order)
    cm_df = pd.DataFrame(cm, index=label_order, columns=label_order)
    display(cm_df)

# Confusion matrix heatmap for visuals
plt.figure(figsize=(6, 5))
plt.imshow(cm, interpolation='nearest')
plt.title(f'Confusion matrix: {name}')
plt.colorbar()
tick_marks = range(len(label_order))
plt.xticks(tick_marks, label_order, rotation=90)
plt.yticks(tick_marks, label_order)
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.tight_layout()
plt.show()
```


Logistic\_L2 holdout accuracy: 0.8683

	Insufficient_Weight	Normal_Weight	Obesity_Type_I	Obesity_Type_II
Insufficient_Weight	479	25	0	0
Normal_Weight	58	505	2	0
Obesity_Type_I	1	0	492	0
Obesity_Type_II	0	0	22	62
Obesity_Type_III	0	0	0	0
Overweight_Level_I	1	46	15	0
Overweight_Level_II	0	3	79	0




LDA holdout accuracy: 0.8230

	Insufficient_Weight	Normal_Weight	Obesity_Type_I	Obesity_Type_II
Insufficient_Weight	470	33	0	5
Normal_Weight	101	442	1	6
Obesity_Type_I	1	1	456	5
Obesity_Type_II	0	0	29	6
Obesity_Type_III	0	0	1	1
Overweight_Level_I	6	74	16	5
Overweight_Level_II	0	9	82	6




GaussianNB holdout accuracy: 0.5860

	Insufficient_Weight	Normal_Weight	Obesity_Type_I	Obesity_Type_II
Insufficient_Weight	377	25	89	6
Normal_Weight	257	123	82	6
Obesity_Type_I	1	17	239	27
Obesity_Type_II	0	2	5	62
Obesity_Type_III	0	1	0	0
Overweight_Level_I	42	53	113	10
Overweight_Level_II	2	29	77	22



SVM\_RBF holdout accuracy: 0.8810

	Insufficient_Weight	Normal_Weight	Obesity_Type_I	Obesity_Type_II
Insufficient_Weight	472	32	0	0
Normal_Weight	47	510	1	0
Obesity_Type_I	1	0	513	0
Obesity_Type_II	0	0	19	63
Obesity_Type_III	0	0	2	0
Overweight_Level_I	2	47	18	0
Overweight_Level_II	0	6	47	0



## 7. Holdout accuracy comparison table

```
In [59]: holdout_df = pd.DataFrame(validation_results).sort_values(by='Holdout_Accuracy')
display(holdout_df)
```

	Model	Holdout_Accuracy
0	SVM_RBF	0.881021
1	Logistic_L2	0.868256
2	LDA	0.822977
3	GaussianNB	0.585983

```
In [60]: # Model Assumptions – Normality of Predictors
# Numeric predictors were evaluated for distributional validity using Q–Q plots.
# Shapiro–Wilk normality testing (n=5000, seed=42). Subsampled diagnostics can
# preserve approximate Gaussian structure, validating model assumption integrity
# that include parametric inference (LDA, logistic-L2).


import scipy.stats as stats
import matplotlib.pyplot as plt

# Q–Q plots for numeric features
for col in numeric_cols:
    plt.figure()
    stats.probplot(train[col].dropna().sample(5000, random_state=42), dist="norm")
    plt.title(f"Q–Q plot for {col}")
    plt.tight_layout()
    plt.show()


# Shapiro–Wilk test
shapiro_results = []
for col in numeric_cols:
    stat, p = stats.shapiro(train[col].dropna().sample(5000, random_state=42))
    shapiro_results.append({"Feature": col, "Shapiro_p_value": p})



pd.DataFrame(shapiro_results).sort_values(by="Shapiro_p_value")
```

Q-Q plot for Age




Q-Q plot for Height







Q-Q plot for NCP




Q-Q plot for CH2O



Q-Q plot for FAF



Q-Q plot for TUE



Out[60]:

	Feature	Shapiro_p_value
4	NCP	7.699289e-75
3	FCVC	8.958148e-63
7	TUE	1.324496e-55
0	Age	6.491899e-52
6	FAF	1.177001e-49
5	CH2O	8.770947e-48
2	Weight	5.619416e-33
1	Height	3.346691e-18

## 8. Train final models on full data and generate prediction files

In [61]:

```
final_predictions = {}
X_test_features = test.drop(columns=['id'])

for name, clf in base_models.items():
    pipe = Pipeline([('pre', preprocess), ('clf', clf)])
    pipe.fit(X, y)
    preds = pipe.predict(X_test_features)
    final_predictions[name] = preds
    out_df = pd.DataFrame({'id': test['id'], target: preds})
    out_name = f'submission_{name}.csv'
    out_df.to_csv(out_name, index=False)
    print(f'Saved {out_name}'')
```

```
Saved submission_Logistic_L2.csv
Saved submission_LDA.csv
Saved submission_GaussianNB.csv
Saved submission_SVM_RBF.csv
```

## 1. Data loading and initial structure

In [62]:

```
import pandas as pd
import numpy as np
from pathlib import Path

# Expect train.csv, test.csv, and sample_submission.csv to be in the same folder
base_path = Path('.')
train_path = base_path / 'train.csv'
test_path = base_path / 'test.csv'
sample_path = base_path / 'sample_submission.csv'

train = pd.read_csv(train_path)
test = pd.read_csv(test_path)
sample = pd.read_csv(sample_path)
```

```

target = 'NObeyesdad'
X = train.drop(columns=[target, 'id'])
y = train[target]

X.head()

```

Out[62]:

	Gender	Age	Height	Weight	family_history_with_overweight	FAVC	
0	Male	24.443011	1.699998	81.669950		yes	yes 2.
1	Female	18.000000	1.560000	57.000000		yes	yes 2.
2	Female	18.000000	1.711460	50.165754		yes	yes 1.
3	Female	20.952737	1.710730	131.274851		yes	yes 3.
4	Male	31.641081	1.914186	93.798055		yes	yes 2.

## 2. Feature types and preprocessing pipeline

In [63]:

```

from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.pipeline import Pipeline

numeric_cols = X.select_dtypes(include=[np.number]).columns.tolist()
categorical_cols = X.select_dtypes(exclude=[np.number]).columns.tolist()
print('Numeric columns:', numeric_cols)
print('Categorical columns:', categorical_cols)

preprocess = ColumnTransformer([
    ('num', StandardScaler(), numeric_cols),
    ('cat', OneHotEncoder(handle_unknown='ignore', sparse_output=False), cat)])

```

Numeric columns: ['Age', 'Height', 'Weight', 'FCVC', 'NCP', 'CH20', 'FAF', 'TUE']  
Categorical columns: ['Gender', 'family\_history\_with\_overweight', 'FAVC', 'CAEC', 'SMOKE', 'SCC', 'CALC', 'MTRANS']

## 3. Class distribution and numeric correlations

In [64]:

```


import matplotlib.pyplot as plt


# Class distribution
plt.figure()
train[target].value_counts().plot(kind='bar')
plt.title('Class distribution for NObeyesdad')
plt.xlabel('Class')
plt.ylabel('Count')
plt.tight_layout()
plt.show()

# Correlation heatmap for numeric features

```

```
corr = train[numeric_cols].corr()
plt.figure(figsize=(8, 6))
im = plt.imshow(corr, interpolation='nearest')
plt.title('Numeric feature correlation matrix')
plt.colorbar(im, fraction=0.046, pad=0.04)
plt.xticks(range(len(numeric_cols)), numeric_cols, rotation=90)
plt.yticks(range(len(numeric_cols)), numeric_cols)
plt.tight_layout()
plt.show()
```





## 4. Cross-validated model performance (four classifiers)

```
In [65]: from sklearn.model_selection import StratifiedKFold, cross_val_score
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

def train_and_eval(name, model):
    pipe = Pipeline([('pre', preprocess), ('clf', model)])
    scores = cross_val_score(pipe, X, y, cv=cv, scoring='accuracy')
    print(f"{name}: {scores.mean():.4f} mean ± {scores.std():.4f} sd")
    pipe.fit(X, y)
    return pipe

# 1. Multinomial logistic regression with L2 regularization
logit_model = train_and_eval('Multinomial logistic regression (L2)', Logisti
```

```

# 2. Linear Discriminant Analysis
lda_model = train_and_eval('Linear Discriminant Analysis', LinearDiscriminat

# 3. Gaussian Naive Bayes
nb_model = train_and_eval('Gaussian Naive Bayes', GaussianNB())

# 4. Support Vector Machine with RBF kernel
svm_model = train_and_eval('RBF-kernel Support Vector Machine', SVC(kernel='

Multinomial logistic regression (L2): 0.8620 mean ± 0.0046 sd
Linear Discriminant Analysis: 0.8207 mean ± 0.0032 sd
Gaussian Naive Bayes: 0.5874 mean ± 0.0053 sd
RBF-kernel Support Vector Machine: 0.8788 mean ± 0.0058 sd

```

In [66]:

```

## Model Diagnostics – Residual Stability & Homoscedasticity

#####
To assess model partition stability, residual spread for the top-performing
#####

# Scale-Location plot function
def scale_location_plot(name, model_pipe, X_data, y_data):
    preds = model_pipe.predict(X_data)
    resid = y_data.astype('category').cat.codes - preds.astype('category').c
    std_resid = np.sqrt(np.abs(resid))

    plt.figure()
    plt.scatter(preds.astype('category').cat.codes, std_resid)
    plt.title(f"Scale-Location plot for {name}")
    plt.xlabel("Fitted Values")
    plt.ylabel("sqrt|Residuals|")
    plt.tight_layout()
    plt.show()

scale_location_plot("SVM-RBF", svm_model, X_val, y_val)
scale_location_plot("Logistic-L2", logit_model, X_val, y_val)

```

---

**TypeError** Traceback (most recent call last)

Cell In[66], line 21

```

18     plt.tight_layout()
19     plt.show()
--> 21 scale_location_plot("SVM-RBF", svm_model, X_val, y_val)
22 scale_location_plot("Logistic-L2", logit_model, X_val, y_val)

```

Cell In[66], line 10, in scale\_location\_plot(name, model\_pipe, X\_data, y\_dat
a)

```

8 def scale_location_plot(name, model_pipe, X_data, y_data):
9     preds = model_pipe.predict(X_data)
--> 10    resid = y_data.astype('category').cat.codes - preds.astype('cate
gory').cat.codes
11    std_resid = np.sqrt(np.abs(resid))
13    plt.figure()

```

**TypeError**: data type 'category' not understood

## 5. Validation split and confusion matrices

```
In [ ]: from sklearn.model_selection import train_test_split

X_train, X_val, y_train, y_val = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)

models_for_validation = [
    ('Multinomial logistic regression (L2)', LogisticRegression(penalty='l2'),
     ('Linear Discriminant Analysis', LinearDiscriminantAnalysis()),
     ('Gaussian Naive Bayes', GaussianNB()),
     ('RBF-kernel Support Vector Machine', SVC(kernel='rbf', C=1, gamma='scal
]

for name, clf in models_for_validation:
    pipe = Pipeline([('pre', preprocess), ('clf', clf)])
    pipe.fit(X_train, y_train)
    preds = pipe.predict(X_val)
    acc = accuracy_score(y_val, preds)
    print(f"\n{name} validation accuracy: {acc:.4f}")
    cm = confusion_matrix(y_val, preds, labels=sorted(y.unique()))
    cm_df = pd.DataFrame(cm, index=sorted(y.unique()), columns=sorted(y.unique))
    print(cm_df)
```

## 6. Generate prediction files for all four models

```
In [ ]: # Fit each model on the full training data and generate prediction files
logit_pipe = Pipeline([('pre', preprocess), ('clf', LogisticRegression(penalty='l2'))])
lda_pipe = Pipeline([('pre', preprocess), ('clf', LinearDiscriminantAnalysis())])
nb_pipe = Pipeline([('pre', preprocess), ('clf', GaussianNB())])
svm_pipe = Pipeline([('pre', preprocess), ('clf', SVC(kernel='rbf', C=1, gamma='scale'))])

logit_pipe.fit(X, y)
lda_pipe.fit(X, y)
nb_pipe.fit(X, y)
svm_pipe.fit(X, y)

logit_preds = logit_pipe.predict(test.drop(columns=['id']))
lda_preds = lda_pipe.predict(test.drop(columns=['id']))
nb_preds = nb_pipe.predict(test.drop(columns=['id']))
svm_preds = svm_pipe.predict(test.drop(columns=['id']))

pd.DataFrame({'id': test['id'], target: logit_preds}).to_csv('submission_logistic_l2.csv')
pd.DataFrame({'id': test['id'], target: lda_preds}).to_csv('submission_lda.csv')
pd.DataFrame({'id': test['id'], target: nb_preds}).to_csv('submission_gaussian_nb.csv')
pd.DataFrame({'id': test['id'], target: svm_preds}).to_csv('submission_svm_r.csv')

print('Prediction files generated:')
print(' - submission_logistic_l2.csv')
print(' - submission_lda.csv')
```

```
print(' - submission_gaussian_nb.csv')
print(' - submission_svm_rbf.csv')
```

In [ ]: *## Final Model Selection & Inference Justification*

The RBF-kernel Support Vector Machine classifier was selected **as** the final m